

An address may be arbitrarily complex, so long as it yields a value within the bounds of the segment it refers to. (Again, for an absolute file, any address will be accepted.) There can be any number of addresses preceding a command so long as the last one or two are legal for that command. For a command requiring two addresses, it is an error for the second address to be less than the first, or for the addresses to refer to different segments. In debug mode, when pointing at the core file, stack addresses are accepted.

In machine instruction mode, addresses should be specified with care. db will disassemble starting at the address specified, but if the address is not the beginning of a valid instruction, the output is meaningless. To insure valid disassembly, start at the beginning of the text section or at the address of a text-relative symbol.

Command Syntax

Each db command consists of a single character, which may optionally be preceded by one or more addresses typically specifying the inclusive range of bytes in the file that the command is to affect. Addresses are separated by a comma ',' or by a colon ':'. If a colon is used, then dot is set to the location and segment specified by the address to the left of the colon before the rest of the command is interpreted. The command character may be followed by additional parameters. Multiple commands may be entered on a line by separating them with semicolons. Every command needing addresses has default values, so addresses can often be omitted; db complains if an address is referred to that doesn't exist.

In the command descriptions below, addresses in parentheses are the default byte addresses that will be used for the command if none are entered. The parentheses are not to be entered with the command, but are present as a syntactical notation. If a command is described without parentheses, no addresses are required. For a two-address command, supplying only one address will default the second to be identical with the one supplied. It is an error to precede with addresses commands that do not require any. The commands are:

(fp)a - display auto stack frame. Valid only if core is the current file. db keeps a pointer fp to the current stack frame, initialized to the one specified by the panel dumped to the core file. "a" may be preceded by an address which will be interpreted (and memorized) as the address of an auto stack frame.

b(1) - backup one or more stack frames. Valid only if core is the current file. "b" may be followed by a zero, which reinitializes the current stack frame pointer fp to the one in the panel, or by a count of the number of frames to be backed over. "b" and this count may also be followed by "a", which will backup one stack frame and display the auto stack frame values.

c - display the panel, which consists of the saved registers and cause of death. The registers are displayed in hexadecimal and symbolically, if meaningful. Valid only if core is the current file.

(.,.)d - print items that differ between executable and core files.

e file - enter a new modifiable data file, in same mode as before, then memorize filename. This works with -a flag only.

f - display name of modifiable data file.

(.)kx - mark address by memorizing it in x, where x is any lower case letter, suitable for later recall as a pointer with 'x'.

(.,.)l<mode> - look at address specified, i.e., display item in output text representation, but don't change '.' to the address given. If <mode> is present, it changes the mode that was specified in the last 'o' command.

(.)n - display address as symbol plus offset.

o<mode> - change output text representation to <mode>, where <mode> is 'm', indicating disassembly of machine instructions, or any of the combinations [a|h|o|u] [c|s|i|l], for ASCII, hexadecimal, octal, or unsigned display of char, short, int, or long integers. If <mode> is not specified, the mode is set to the default used at db invocation: 'm' if the -a flag was not given, and unsigned short otherwise.

(.,.)p<mode> - print item in output text representation, then change '.' to point at last address specified. If <mode> is present, it changes the mode that was specified in the last 'o' command.

q - quit.

(.)r file - read binary contents of specified file and replace target file from current byte, extending the file as necessary. This is valid only with -a flag. file must not be the <file> being db'ed.

(.,.)s/pattern/new/[g] - for all items in range, expand to output text representation, edit text by changing instances of pattern to new, then update item using edited text pattern. If g suffix is specified, all matches in an item are altered, otherwise only the leftmost is altered.

(.)u - enter update mode, i.e., treat each subsequent line of input as the output text representation of an item and replace the addressed item. Successive items are replaced, extending the file as necessary, until an input line containing only a '.' is encountered. This does not work in 'm' mode. The input must be entered in the current mode, one item per line; otherwise db will complain.

(0,\$)w file - write the specified range of the file being db'ed to file. file must not be the same as <file> being db'ed. This command is valid only with -a flag.

z file - point db at file. If file is omitted, z reports which file, core or executable, is currently pointed at.

(.)<newline> - equivalent to entering ". + size", where size is the number of bytes being displayed in the current mode. This form can be used to view one instruction or item at a time. If a byte address precedes the <newline>, the addressed item will be displayed. Dot will be set to the displayed byte.

!<command> - send all characters on the line to the right of the '!' to the system, to be interpreted as a system command. "!cd" is handled by db directly. Occurrences of the sequence "\f" will each be replaced by the currently remembered file before the transfer to system level is made. Dot is not changed by this command. Available only in Idris/UNIX systems.

= - display the byte number of dot.

On Idris/UNIX systems, if an ASCII DEL is typed, db interrupts its current task, if any, and displays '?'. It then will be ready to accept db commands again. Dot is generally ill-defined at this point.

SEE ALSO

link, rel

BUGS

This utility is currently provided for use only on Idris/UNIX host systems.

NAME

hex - translate object file to ASCII formats

SYNOPSIS

hex -[db## dr h m* r## s tb## +# #] <ofile>

FUNCTION

hex translates executable images produced by link to Intel standard hex format or to Motorola S-record object file format. The executable image is read from <ofile>. If no <ofile> is given, or if the filename given is "-", the file xeq is read.

The flags are:

-db## when processing a standard object file as input, override the object module data bias with ##.

-dr output text records as record type 0x81, data records as record type 0x82, for use with the Digital Research genccmd utility under CP/M-86.

-h don't produce start record for Intel hex files.

-m* insert the string *, instead of <ofile>, into the Motorola S0 record generated when -s is given.

-r## interpret the input file as "raw" binary and not as an object file. Output is produced as described below, in either format, except that the starting address of the module is specified by the long integer ##.

-s produce S-records rather than the default hex format.

-tb## when processing a standard object file as input, override the object module text bias with ##.

+# start output with the #th byte. # should be between 0 and one less than the value specified by the -# flag. -4 +3 produces bytes 3, 7, 11, 15, ...; -2 +0 produces all even bytes; -2 +1 outputs all odd bytes; 0 is the default, which outputs all bytes.

-# output every #th byte. -2 outputs every other byte, -4 every fourth; 1 is the default, which outputs all bytes.

Output is written to STDOUT. When the input file is a standard object file, the load offset of the first record output for its text or data segment is determined as follows. If an offset is explicitly given for that segment (with "-db##" or "-tb##"), it is used. Otherwise, if all bytes of the input file are being output, the appropriate bias is used from the object file header. Otherwise, a load offset of zero is used. Then, the value given with "+#" is added to the chosen offset to obtain the first offset actually output.

Hex Files

A file in Intel hex format consists of the following records, in the order listed:

- 1) a '\$' alone on a line to indicate the end of the (non-existent) symbol table. If -h is specified, this line is omitted.
- 2) data records for the text segment, if any. These represent 32 image bytes per line, possibly terminated by a shorter line.
- 3) data records for the data segment, if any, in the same format as the text segment records.
- 4) an end record specifying the start address.

Data records each begin with a ':' and consist of pairs of hexadecimal digits, each pair representing the numerical value of a byte. The last pair is a checksum such that the numerical sum of all the bytes represented on the line, modulo 256, is zero. The bytes on a data record line are:

- a) the number of image bytes on the line.
- b) two bytes for the load offset of the first image byte. The offset is written more significant byte first so that it reads correctly as a four-digit hexadecimal number.
- c) a zero byte "00".
- d) the image bytes in increasing order of address.
- e) the checksum byte.

An end record also begins with a ':' and is written as digit pairs with a trailing checksum. Its format is:

- a) a zero byte "00".
- b) two bytes for the start address, in the same format as the load offset in a data record. The start address used is the first load offset output for the file.
- c) a one byte "01".
- d) the checksum byte.

S-Record Files

A file in Motorola S-record format is a series of records each containing the following fields:

<S field><count><addr><data bytes><checksum>

All information is represented as pairs of hexadecimal digits, each pair representing the numerical value of a byte.

<S field> determines the interpretation of the remainder of the line; valid S fields are "S0", "S1", "S2", "S8", "S9". <count> indicates the number of bytes represented in the rest of the line, so the total number of characters in the line is <count> * 2 + 4.

<addr> indicates the byte address of the first data byte in the data field. S0 records have two zero bytes as their address field; S1 and S2 records have <addr> fields of two and three bytes in length, respectively; S9 and S8 records have <addr> fields of two and three bytes in length, respectively, and contain no data bytes. <addr> is represented most significant byte first.

The S0 record contains the name of the input file, formatted as data bytes. If input was from xeq, XEQ is used as the name. S1 and S2 records represent text or data segment bytes to be loaded. They normally contain 32 image bytes, output in increasing order of address; the last record of each segment may be shorter. The text segment is output first, followed by the data segment. S9 records contain only a two-byte start address in their <addr> field; S8 records contain a three-byte address. The start address used is the first load offset output for the file.

<checksum> is a single byte value such that the numerical sum of all the bytes represented on the line (except the S field), taken modulo 256, is 255 (0xFF).

RETURNS

hex returns success if no error messages are printed, that is, if all records make sense and all reads and writes succeed; otherwise it reports failure.

EXAMPLE

The file hello.c, consisting of:

```
char *p {"hello world"};
```

when compiled produces the following Intel hex file:

```
% hex hello.o
$
:0E00000020068656C6C6F20776F726C640094
:00000001FF
```

and the following Motorola S-records:

```
% hex -s hello.o
S00A000068656C6C6F2E6F44
S111000020068656C6C6F20776F726C640090
S9030000FC
```

SEE ALSO

link, rel

NAME

lib - maintain libraries

SYNOPSIS

lib <lfile> -[c d i p r t v3 v6 v7 v x] <files>

FUNCTION

lib provides all the functions necessary to build and maintain object module libraries in either standard library format or those used by UNIX/System III, UNIX/V6, or UNIX/V7. lib can also be used to collect arbitrary files into one bucket. <lfile> is the name of an existing library file or, in the case of replace or create operations, the name of the library to be constructed.

The flags are:

- c create a library containing <files>, in the order specified. Any existing <lfile> is removed before the new one is created.
- d delete from the library the zero or more files in <files>.
- i take <files> from STDIN instead of the command line. Any <files> on the command line are ignored.
- p print named files (do a -x to STDOUT). If no files are named, all files are extracted.
- r in an existing library, replace the zero or more files in <files>; if no library <lfile> exists, create a library containing <files>, in the order specified. The files in <files> not present in the library are appended to it, in the order specified.
- t list the files in the library in the order they occur. If <files> are given, then only the files named and present are listed.
- v3 create the output library in UNIX/System III format, VAX-11 byte-order. This flag is meaningful with -c or -r; it is mandatory for -t and -x (UNIX/V7 format is assumed if this flag is not specified).
- v6 create the output library in UNIX/V6 format. Meaningful only with -c or -r.
- v7 create the output library in UNIX/V7 format, PDP-11 byte-order. Meaningful only with -c or -r.
- v be verbose. The name of each object file operated on is output, preceded by a code indicating its status: "c" for files retained unmodified, "d" for those deleted, "r" for those replaced, and "a" for those appended to <lfile>. If -v is used in conjunction with -t, each object file is listed followed by its length in bytes.

- x extract the files in <files> that are present in the library into discrete files with the same names. If no <files> are given, all files in the library are extracted.

At most one of the flags -[c d p r t x] may be given; if none is given, -t is assumed. Similarly, at most one of the flags -[v3 v6 v7] should be present; if none is present, standard library file format is assumed when a new library is created by -r or -c. An existing library is processed in its proper mode, except for UNIX/III libraries, for which the flag -v3 must be specified.

The standard library format consists of a two-byte header having the value 0177565, written less significant byte first, followed by zero or more entries. Each entry consists of a fourteen-byte name, NUL padded, followed by a two-byte unsigned file length, also stored less significant byte first, followed by the file contents proper. If a name begins with a NUL byte, it is taken as the end of the library file.

Note that this differs in several small ways from UNIX/V6 format, which has a header of 0177555, an eight-byte name, six bytes of miscellaneous UNIX-specific file attributes, and a two-byte file length. Moreover, a file whose length is odd is followed by a NUL padding byte in the UNIX format, while no padding is used in standard library format.

UNIX/III and UNIX/V7 formats are characterized by a header of 0177545, a fourteen-byte name, eight bytes of UNIX-specific file attributes, and a four-byte length. The length is in VAX-11 native byte order for UNIX/III or in PDP-11 native byte order for UNIX/V7. Odd length files are also padded to even.

RETURNS

lib returns success if no problems are encountered, else failure. After most failures, an error message is printed to STDERR and the library file is not modified. Output from the -t flag, and verbose remarks, are written to STDOUT.

EXAMPLE

To build a library and check its contents:

```
% lib clib -r one.o two.o three.o  
% lib clib -tv
```

SEE ALSO

link, lord, rel

BUGS

If all modules are deleted from a library, a vestigial file remains.

Modifying UNIX/III, UNIX/V6, or UNIX/V7 format files causes all attributes of all file entries to be zeroed.

- 3 -

lib doesn't check for files too large to be properly represented in a library (> 65,534 bytes).

NAME

link - combine object files

SYNOPSIS

```
link -[a bb## b## c db## dr# d eb* ed* et* h i l*^ o* r sb* sd* st*
      tb## tf# t u*^ x#] <files>
```

FUNCTION

link combines relocatable object files in standard format for any target machine, selectively loading from libraries of such files made with lib, to create an executable image for operation under Idris, or for stand alone execution, or for input to other binary reformatters.

The flags are:

- a make all relocation items and all symbols absolute. Used to prerelocate code that will be linked in elsewhere, and is not to be re-relocated.
- bb## relocate the bss (block started by symbol) segment to start at ##. By default, the bss segment is relocated relative to the end of the data segment. Once -bb## has been specified, the resulting output file is unsuitable for input to another invocation of link.
- b## set the stack plus heap size in the output module header to ##. Idris takes this value, if non-zero, as the minimum number of bytes to reserve at execution time for stack and data area growth. If the value is odd, Idris will execute the program with a smaller heap plus stack size, so long as some irreducible minimum space can still be provided by using all of available memory.
- c suppress code output (.text and .data), and make all symbols absolute. Used to make a module that only defines symbol values, for specifying addresses in shared libraries, etc.
- db## set data bias to the long integer ##. Default is end of text section, rounded up to required storage boundary for the target machine.
- dr# round data bias up to ensure that there are at least # low-order binary zeros in its value. Ignored if -db## is specified.
- d do not define bss symbols, and do not complain about undefined symbols. Used for partial links, i.e., if the output module is to be input to link.
- eb* if the symbol * is referenced, make it equal to the first unused location past the bss area.
- ed* if the symbol * is referenced, make it equal to the first unused location past the initialized data area.
- et* if the symbol * is referenced, make it equal to the first unused location past the text area.

- h suppress header in output file. This should be specified only for stand alone modules such as bootstraps.
- i take <files> from STDIN instead of the command line. Any <files> on the command line are ignored.
- l* append library name to the end of the list of files to be linked, where the library name is formed by appending * to "/lib/lib". Thus "-lc.11" produces "/lib/libc.11". Up to ten such names may be specified as flags; they are appended in the order specified.
- o* write output module to file *. Default is xeq.
- r suppress relocation bits. This should not be specified if the output module is to be input to link or executed under a version of Idris that relocates commands on startup.
- sb* if the symbol * is referenced, make it equal to the size of the bss area.
- sd* if the symbol * is referenced, make it equal to the size of the data area.
- st* if the symbol * is referenced, make it equal to the size of the text area.
- tb## set text bias to the long integer ##. Default is location zero.
- tf# round text size up by appending NUL bytes to the text section, to ensure that there are at least # low-order binary zeroes in the resulting value. Default is zero, i.e. no padding.
- t suppress symbol table. This should not be specified if the output module is to be input to link.
- u* enter the symbol * into the symbol table as an undefined public reference, usually to force loading of selected modules from a library.
- x# specify placement in the output module of the input .text and .data sections. The 2-weighted bit of # routes .text input, the 1-weighted bit, .data. If either bit is set, the corresponding sections are put in the text segment output; otherwise, they go in the data segment. The default value used, predictably, is 2.

The bss section is always assumed to follow the data section in all input files. Unless -bb## is given, the bss section will be made to follow the data section in the output file, as well. It is perfectly permissible for text and data sections to overlap, as far as link is concerned; the target machine may or may not make sense of this situation (as with separate instruction and data spaces).

The specified <files> are linked in order; if a file is in library format, it is searched once from start to end. Only those library modules

are included which define public symbols for which there are currently outstanding unsatisfied references. Hence, libraries must be carefully ordered, or rescanned, to ensure that all references are resolved. By special dispensation, flags of the form "-l*" may be interspersed among <files>. These call for the corresponding libraries to be searched at the points specified in the list of files. No space may occur after the "-l", in this usage.

File Format

A relocatable object image consists of a header followed by a text segment, a data segment, the symbol table, and relocation information.

The header consists of an identification byte 0x99, a configuration byte, a short unsigned int containing the number of symbol table bytes, and six unsigned ints giving: the number of bytes of object code defined by the text segment, the number of bytes of object code defined by the data segment, the number of bytes needed by the bss segment, the number of bytes needed for stack plus heap, the text segment offset, and the data segment offset.

Byte order and size of all ints in the header are determined by the configuration byte.

The configuration byte contains all information needed to fully represent the header and remaining information in the file. Its value val defines the following fields: $((val \& 07) \ll 1) + 1$ is the number of characters in the symbol table name field, so that values [0, 8) provide for odd lengths in the range [1, 15]. If $(val \& 010)$ then ints are four bytes; otherwise they are two bytes. If $(val \& 020)$ then ints in the data segment are represented least significant byte first, otherwise, most significant byte first; byte order is assumed to be purely ascending or purely descending. If $(val \& 040)$ then even byte boundaries are enforced by the hardware. If $(val \& 0100)$ then the text segment byte order is reversed from that of the data segment; otherwise text segment byte order is the same. If $(val \& 0200)$ no relocation information is present in this file.

The text segment is relocated relative to the text segment offset given in the header (usually zero), while the data segment is relocated relative to the data segment offset (usually the end of the text segment). Unless $-bb\#$ is given, the bss segment is relocated relative to the end of the data segment.

Relocation information consists of two successive byte streams, one for the text segment and one for the data segment, each terminated by a zero control byte. Control bytes in the range [1, 31] cause that many bytes in the corresponding segment to be skipped; bytes in the range [32, 63] skip 32 bytes, plus 256 times the control byte minus 32, plus the number of bytes specified by the relocation byte following.

All other control bytes control relocation of the next short or long int in the corresponding segment. If the 1-weighted bit is set in such a control byte, then a change in load bias must be subtracted from the int.

The 2-weighted bit is set if a long int is being relocated instead of a short int. The value of the control byte right-shifted two places, minus 16, constitutes a "symbol code".

A symbol code of 47 is replaced by a code obtained from the byte or bytes following in the relocation stream. If the next byte is less than 128, then the symbol code is its value plus 47. Otherwise, the code is that byte minus 128, times 256, plus 175 plus the value of the next relocation byte after that one.

A symbol code of zero calls for no further relocation; 1 means that a change in text bias must be added to the item (short or long int); 2 means that a change in data bias must be added; 3 means that a change in bss bias must be added. Other symbol codes call for the value of the symbol table entry indexed by the symbol code minus 4 to be added to the item.

Each symbol table entry consists of a value int, a flag byte, and a name padded with trailing NULs. Meaningful flag values are 0 for undefined, 4 for defined absolute, 5 for defined text relative, 6 for defined data relative, and 7 for defined bss relative. To this is added 010 if the symbol is to be globally known. If a symbol still undefined after linking has a non-zero value, link assigns the symbol a unique area, in the bss segment, whose length is specified by the value, and considers the symbol defined. This occurs only if -d has not been given.

RETURNS

link returns success if no error messages are printed to STDOUT, that is, if no undefined symbols remain and if all reads and writes succeed; otherwise it returns failure.

EXAMPLE

To load the C program echo.o under Idris/S11:

```
% link -lc.11 -t /lib/Crts.o echo.o
```

with separate I/D spaces, for UNIX:

```
% link -lc.11 -rt -db0 /lib/Crts.o echo.o; taout
```

or with read only text section for UNIX:

```
% link -lc.11 -rt -dr13 /lib/Crts.o echo.o; taout
```

And to load the 8080 version of echo under CP/M:

```
A:link -hrt -tb0x0100 a:chdr.o echo.o a:clib.a a:mlib.a
```

SEE ALSO

hex, lib, lord, rel

NAME

lord - order libraries

SYNOPSIS

lord -[c* d*^ i r*^ s]

FUNCTION

lord reads in a list of module names, with associated interdependencies, from STDIN, and outputs to STDOUT a topologically sorted list of module names such that, if at all possible, no module depends on an earlier module in the list. Each module is introduced by a line containing its name followed by a colon. Subsequent lines are interpreted as either:

defs - things defined by the module,

refs - things referred to by the module, or

other stuff - other stuff.

Refs and defs have the syntax given by one or more formats entered as flags on the command line. Each character of the format must match the corresponding character at the beginning of an input line; a ? will match any character except newline. If all the characters of the format match, then the rest of the input line is taken as a ref or def name. Thus, the format flag "-d0x????D" would identify as a valid def any line beginning with "0x", four arbitrary characters and a "D", so that the input line "0x3ff0D _inbuf" would be taken as a def named "_inbuf".

The flags are:

-c* prepend the string * to the output stream. Implies -s. Each module name is output preceded by a space; the output stream is terminated with a newline. Hence, lord can be used to build a command line.

-d* use the string * as a format for defs.

-i ignore other stuff. Default is to complain about any line not recognizable as a def or ref.

-r* use the string * as a format for refs.

-s suppress output of defs and refs; output only module names in order.

Up to ten formats may be input for defs, and up to ten for refs.

If no -d flags are given, lord uses the default def formats: "0x??????B", "0x??????D", "0x??????T", "0x????B", "0x????D", "0x????T". If no -r flags are given, lord uses the default ref formats: "0x??????U" and "0x????U". These are compatible with the default output of rel.

If there are circular dependencies among the modules, lord writes the name of the file that begins the unsorted list to STDERR, followed by the message "not completely sorted". In general, rearrangements are made only

lord

- 2 -

lord

when necessary, so an ordered set of modules should pass through lord unchanged.

RETURNS

lord returns success if no error messages are printed, otherwise failure.

EXAMPLE

To create an ordered library of object modules under Idris:

```
% rel *.o | lord -c"lib libx.a -c" | sh
```

To order a set of objects using UNIX nm:

```
% nm *.o > nmlist
% lord < nmlist -c"ar r libx.a" | \
-d"?????T" -d"??????D" -d"?????B" -r"?????U" | sh
```

SEE ALSO

lib, rel

NAME

p1 - parse C programs

SYNOPSIS

p1 -[a b# c e l m n# o# r# u] <file>

FUNCTION

p1 is the parsing pass of the C compiler. It accepts a sequential file of lexemes from the preprocessor pp and writes a sequential file of flow graphs and parse trees, suitable for input to a machine-dependent code generator p2. The operation of p1 is largely independent of any target machine. The flag options are:

- a compile code for machines with separate address and data registers. This flag implies -r6 (because currently used only in conjunction with the MC68000 code generator).
- b# enforce storage boundaries according to #, which is reduced modulo 4. A bound of 0 leaves no holes in structures or auto allocations; a bound of 1 (default) requires short, int and longer data to begin on an even bound; a bound of 2 is the same as 1, except that 4-8 byte data are forced to a multiple of four byte boundary; a bound of 3 is the same as 2, except that 8 byte data (doubles) are forced to a multiple of eight byte boundary.
- c ignore case distinctions in testing external identifiers for equality, and map all names to lowercase on output. By default, case distinctions matter.
- e don't force loading of extern references that are declared but never defined or used in an expression. Default is to load all externs declared.
- l take integers and pointers to be 4 bytes long. Default is 2 bytes.
- m treat each struct/union as a separate name space, and require x.m to have x a structure with m one of its members.
- n# ignore characters after the first # in testing external identifiers for equality. Default is 7; maximum is 8, except that values up to 32 are permitted for the VAX-11 code generator only.
- o# write the output to the file # and write error messages to STDOUT. Default is STDOUT for output and STDERR for error messages.
- r# assign no more than # variables to registers at any one time, where # is reduced modulo 7. Default is 3 register variables; values above 3 are currently acceptable only for the MC68000 code generator (3 data registers + 3 address registers maximum), and the VAX-11 code generator (6 registers maximum).
- u take "string" as array of unsigned char, not array of char.

If <file> is present, it is used as the input file instead of the default STDIN. On many systems (other than Idris/UNIX), the -o option and <file> are mandatory because STDIN and STDOUT are interpreted as text files, and hence become corrupted.

EXAMPLE

p1 is usually sandwiched between pp and some version of p2, as in:

```
pp -x -o temp1 file.c  
p1 -o temp2 temp1  
p2.11 -o file.s temp2
```

SEE ALSO

pp

BUGS

p1 can be rather cavalier about semicolons.

NAME

p2.11 - generate code for PDP-11 C programs

SYNOPSIS

p2.11 -[ck e f n* o* p x#] <file>

FUNCTION

p2.11 is the code generating pass of the C compiler. It accepts a sequential file of flow graphs and parse trees from p1 and writes a sequential file of assembly language statements, suitable for input to a PDP-11 assembler. Two versions of p2.11 are available for the PDP-11: one that generates code compatible with the DEC MACRO-11 assemblers, and one that is compatible with Idris and UNIX assemblers.

As much as possible, the compiler generates free-standing code; but for those operations which cannot be done compactly, it generates inline calls to a set of machine-dependent runtime library routines. The PDP-11 runtime library is documented in Section IV of this manual.

The flags are:

- ck enable stack overflow checking.
- e generate code for non-EIS machine, e.g., an 11/04 or an 11/03 without the FIS chip. Default is to use ash, ashc, mul, div, and sxt instructions as needed.
- f generate code for machine with hardware Floating Point Processor. Default is inline calls to runtime floating routines.
- n* use the name * as the module name in the assembler output. Default is to use the first defined external encountered.
- p emit profiler calls on entry to each function. Usable with UNIX profiler.
- o* write the output to the file * and write error messages to STDOUT. Default is STDOUT for output and STDERR for error messages.
- x# map the three virtual sections, for Functions (04), Literals (02), and Variables (01), to the two physical sections, Code (bit is a one) and Data (bit is a zero). Thus, "-x4" is for separate I/D space, "-x6" is for ROM/RAM code, and "-x7" is for compiling tables into ROM. Default is 4.

Note that code produced with the -f option is not compatible with that produced without the -f option, since one version uses the hardware registers in the FPP to return floating function values and the other uses software registers provided by the library. The portable library is ordinarily provided with substitute routines for use with hardware FPP.

There is no problem in mixing code produced with and without the -e flag.

If <file> is present, it is used as the input file instead of the default STDIN. On many systems (other than Idris/UNIX), <file> is mandatory, because STDIN is interpreted as a text file, and hence becomes corrupted.

Files output from p1 for use with the PDP-11 code generator must be generated with: no -l flag, since pointers are short; no -b# flag, since only the default value of "-b1" is acceptable; and no -a flag, since registers are interchangeable. For use with the MACRO-11 assembler, p1 should be run with the flags "-cn6" to check externals properly; the default "-n7" is suitable for UNIX; "-n8" is preferred for the Idris assembler.

Wherever possible, labels in the emitted code each contain a comment which gives the source line from which the code immediately following obtains, along with a running count of the number of words of code produced for a given function body.

EXAMPLE

p2 usually follows pp and p1, as follows:

```
pp -o temp1 -x file.c
p1 -o temp2 -n8 temp1
p2.11 -o file.s temp2
```

SEE ALSO

as.11, p1

BUGS

Stack overflow checking is only approximate, since a calculation of the exact stack high water mark is not attempted.

NAME

pp - preprocess defines and includes

SYNOPSIS

pp -[c d*^ i* o* p? s? x 6] <files>

FUNCTION

pp is the preprocessor used by the C compiler, to perform #define, #include, and other functions signalled by a #, before actual compilation begins. It can be used to advantage, however, with most language processors. The flag options are:

- c don't strip out /* comments */ nor continue lines that end with \.
- d# where # has the form name=def, define name with the definition string def before reading the input; if =def is omitted, the definition is taken as "1". The name and def must be in the same argument, i.e., no blanks are permitted unless the argument is quoted. Up to ten definitions may be entered in this fashion.
- i# change the prefix used with #include <filename> from the default "" to the string *. Multiple prefixes to be tried in order may be specified, separated by the character '|'.
- o# write the output to the file * and write error messages to STDOUT. Default is STDOUT for output and STDERR for error messages. On many systems (other than Idris/UNIX), the -o option is mandatory with -x because STDOUT is interpreted as a text file, and hence becomes corrupted.
- p? change the preprocessor control character from '#' to the character ?.
- s? change the secondary preprocessor control character from '@' to the character ?.
- x put out lexemes for input to the C compiler p1, not lines of text.
- 6 put out extra newlines and/or SOH ('\1') codes to keep source line numbers correct for UNIX/V6 compiler or ptc.

pp processes the named files, or STDIN if none are given, in the order specified, writing the resultant text to STDOUT. Preprocessor actions are described in detail in Section I of the C Programmers' Manual.

The presence of a secondary preprocessor control character permits two levels of parameterization. For instance, the invocation

pp -c -p@

will expand define and ifdef conditionals, leaving all # commands and comments intact; invoking pp with no arguments would expand both @ and # commands. The flag -s# would effectively disable the secondary control character.

EXAMPLE

The standard style for writing C programs is;

```
/* name of program
 */
#include <std.h>

#define MAXN    100

COUNT things[MAXN];
etc.
```

The use of uppercase only identifiers is not required by pp, but is strongly recommended to distinguish parameters from normal program identifiers and keywords.

SEE ALSO

p1, ptc

BUGS

Unbalanced quotes ' or " may not occur in a line, even in the absence of the -x flag. Floating constants longer than 38 digits may compile incorrectly, on some host machines.

NAME

prof - produce execution profile

SYNOPSIS

prof -[a f +l n r s t +z] [<ofile>] [<pfiles>]

FUNCTION

prof correlates one or more files of statistics recorded during the execution of a program with the symbol table of the corresponding executable file, and produces a report profiling the run. The statistics are contained in "profile" files of standard format. The <ofile> given on the command line is the program whose execution produced the statistics.

Flags are:

- a sort output in increasing order of address.
- f sort output in decreasing order of frequency of calls to each function.
- +l include local symbols in output. By default, only global symbols are included.
- n sort output in increasing order of symbol name.
- r reverse the default sense of the output sort.
- s suppress the default heading.
- t sort output in decreasing order of time spent in each function. This is the default sort key.
- +z include symbols in output for which no entries or execution time were recorded. By default, these are excluded.

If no <ofile> is given, "xeq" is used as the object filename. If no <pfiles> are given, prof expects profiling data to come from the file "profil". However, if any filenames appear after <ofile> on the command line, they will be read for profiling data instead. If more than one file is given, prof will accumulate data from each in turn, and output a single report giving the totals from all data files read. Thus the results of several instrumented runs of a given program may be merged to reduce statistical error. Naturally, if <pfiles> are given <ofile> must be given as well.

Profiling involves two statistics: counting the calls made to each function in a program, and recording what portion of execution time is spent inside each function. In its report, prof outputs one line for each text-relative symbol in <file> (subject to +l and +z), showing the symbol name, the percentage of total execution time spent in the region between this symbol and the next higher one, time in this region in milliseconds, the number of calls made directly to this symbol's address, and the average amount of time recorded per call. Also, prof treats the final location available for function entry counts as an overflow location, in which

calls are counted to all functions that could not be given a location of their own. If this count is non-zero, prof outputs it, with the name "other syms".

Idris records execution time by periodically examining the PC of the executing program, using its value to index an array whose elements correspond to the permissible range of the PC. The selected array element is then incremented. Function entry counts are maintained by calls to a counting routine, which every Whitesmiths C compiler can be made to output by specifying the flag -p to the code generator. A function must be thus instrumented for an entry count to be maintained for it. The counting routine itself is specified at link time, as part of the runtime startup code.

RETURNS

prof returns success if all input files are readable and consistent with their expected format.

SEE ALSO

The profile file format description is in Section III of the Idris Programmers' Manual, while the runtime initialization code used to profile under Idris is documented along with the Idris system interface, in Section III.a of this manual. The portable profiling setup functions are described in Section IV of the Idris Programmers' Manual, and the machine-dependent function entry counting routine is described in Section IV of the current manual.

BUGS

Sampling errors may occur unless the scaling factor used to record the PC location is no larger than the smallest boundary on which an instruction may begin; that is, a given array element may overlap function bodies, and a PC caught in one function may be treated as if it were caught in another. prof does what it can to compensate, by dividing the value of an overlapping array element between the functions involved.

This utility is currently provided for use only on Idris host systems.

NAME

ptc - Pascal to C translator

SYNOPSIS

ptc -[c f k m# n# o* r s#] <infile>

FUNCTION

ptc is a program that accepts as input lines of Pascal text and produces as output a corresponding C program which is acceptable to the Whitesmiths, Ltd. C compiler. If <infile> is present, it is taken as the Pascal program to translate; otherwise input is taken from STDIN.

The flags are:

- c pass comments through to the C program.
- f set the precision for reals to single precision (float). Default is double.
- k permit pointer types to be defined using type identifiers from outer blocks. Default is ISO standard, i.e., the type pointed to must be defined in the same type declaration as the pointer type definition.
- m# make # the number of bits in MAXINT excluding the sign bit, e.g., MAXINT becomes 32767 for -m15, 1 for -m1, etc. Default for MAXINT is 32766 [sic]. Acceptable values for # are in the range [0 , 32]. Declaring MAXINT greater than the size of a pointer (16 or 32) will give unpredictable results. On a target machine with 16-bit pointers, # should be less than 16.
- n# Make # the number of significant characters in external names. Default is 8 character external names.
- o* write the C program to the file * and diagnostics to STDOUT. Default is STDOUT for the C program and STDERR for diagnostics.
- r turn off runtime array bounds checks.
- s# make # the number of bits in the maximum allowable set size, i.e., the size of all sets whose basetype is integer becomes the specified power of two. Acceptable values are in the range [0 , 32]. Default is 8 (256 elements).

The CP/M operating system implementation, on the Intel 8080 and Zilog Z/80, restricts the acceptable value for # to the range [0,16]; for maximum portability, this restriction should be honored.

Identifiers are mapped to uppercase to keep from conflicting with those declared as reserved words in C. Moreover, structure declarations may be produced that contain conflicting field declarations; and declarations are present for library functions that may not be needed. All of these peccadilloes are forgiven by the use of appropriate C compiler options.

RETURNS

ptc returns success if it produces no diagnostics.

BUGS

Complex set expressions can produce very long lines that are truncated by pp.

NAME

rel - examine object files

SYNOPSIS

rel -[d g i o s t u v] <files>

FUNCTION

rel permits inspection of relocatable object files, in standard format, for any target machine. Such files may have been output by an assembler, combined by link, or archived by lib. rel can be used either to check their size and configuration, or to output information from their symbol tables.

The flags are:

- d output all defined symbols in each file, one per line. Each line contains the value of the symbol, a code indicating to what the value is relative, and the symbol name. Values are output as the number of digits needed to represent an integer on the target machine. Relocation codes are: 'T' for text relative, 'D' for data relative, 'B' for bss relative, 'A' for absolute, or '?' for anything rel doesn't recognize. Lowercase letters are used for local symbols, uppercase for globals.
- g print global symbols only.
- i print all global symbols, with the interval in bytes between successive symbols shown in each value field. Implies the flags -[d u v].
- o output symbol values in octal. Default is hexadecimal.
- s display the sizes, in decimal, of the text segment, the data segment, the bss segment, and the space reserved for the runtime stack plus heap, followed by the sum of all the sizes.
- t list type information for this file. For each object file the data output is: the size of an integer on the target machine, the target machine byte order, whether even byte boundaries are enforced by the hardware, whether the text segment byte order is reversed from that of the data segment, and the maximum number of characters permitted in an external name. If the file is a library, then the type of library is output and the information above is output for each module in the library.
- u list all undefined symbols in each file. If -d is also specified, each undefined symbol is listed with the code 'U'. The value of each symbol, if non-zero, is the space to be reserved for it at load time if it is not explicitly defined.
- v sort by value; implies the -d flag above. Symbols of equal value are sorted alphabetically.

If no flags are given, the default is -[d u]; that is, all symbols are listed, sorted in alphabetical order on symbol name. If more than one of

the flags -[d s t u] is selected, then type information is output first, followed by segment sizes, followed by the symbol list specified with -d or -u.

<files> specifies zero or more files, which must be in relocatable format, or standard library format, or UNIX/V6 library format, or UNIX/V7 library format. If more than one file, or a library, is specified, then the name of each separate file or module precedes any information output for it, each name followed by a colon and a newline; if -s is given, a line of totals is also output. If no <files> are specified, or if "-" is encountered on the command line, xeq is used.

RETURNS

rel returns success if no diagnostics are produced, that is, if all reads are successful and all file formats are valid.

EXAMPLE

To obtain a list of all symbols in a module:

```
% rel alloc.o
0x00000074T _alloc
0x00000000U _exit
0x000001feT _free
0x000000beT _malloc
0x00000000U _sbreak
0x00000000U _write
```

SEE ALSO

lib, link, lorc